

ABLS: Working Towards Attribute-Based Log Security and Automated Auditing in the Cloud

Christopher A. Wood
Department of Computer Science
caw4567@rit.edu

ABSTRACT

User-based non-repudiation is an increasingly important property of cloud-based applications. It provides irrefutable evidence that ties system behavior to specific users, thus enabling strict enforcement of organizational security policies. System logs are typically used as the basis for this property. Thus, the effectiveness of system audits based on log files reduces to the problem of maintaining the integrity and confidentiality of log files without sacrificing the usefulness of the data in these log files. In an ideal setting, automated audits would be possible on encrypted log data that defines audit trails. Furthermore, since useful log messages may contain sensitive information, access control for log data should be implemented so as to restrict access to only those parties that need to view it (i.e. generating users, colleagues of generating users, auditors, system administrators, etc). ABAC has been a common technique used to satisfy this requirement.

In this paper we address all of the aforementioned issues with ABLs, an attribute-based logging system designed to support automates audits of encrypted audit trails (log data) based on user-defined security policies. Access to sensitive log information is enforced using ciphertext-policy attribute-based encryption (CP-ABE) with a minimal number of log-related roles, and thus a small number of attributes, to avoid the problem of increasing encryption computational complexity with attribute explosion. We present the preliminary design of ABLs and discuss how audit trails are constructed, automated audit tasks are defined and specified, and how the system may be used in practical applications.

1. INTRODUCTION

User-based non-repudiation is a system security property that provides indisputable evidence linking specific actions to individual users (or entities) that trigger such actions. Cryptographically speaking, non-repudiation requires that the integrity and origin of all data should be provable. In essence, this enables system audits to be conducted that can identify data misuse, and thus, potential security policy violations, by

comparing the contextual information of system events (e.g. source user, time of the event, etc) with all entities authorized to invoke such events. Therefore, treating non-repudiation as a required system quality attribute in the architecture is likely to become a common trend in the commercial, government, and even more specifically, the health-care domain.

System audits typically use log files to determine the “who, what, when, and how” of events that took place during the system’s lifetime. In order to provide accurate information for non-repudiation purposes, it is often necessary to place some amount user-sensitive data in these log files that can be used to trace data back to its origin. As such, logs of events generated by a client that is being served must maintain data confidentiality and integrity should the system be compromised. Historical approaches to the problem of log security are based on tamper-resistant hardware and maintaining continuous secure communication channels between a log aggregator and end user [17]. However, such solutions are not applicable in the context of cloud-based applications.

Recent approaches have relied on combinations of encryption and signature techniques [12]. Symmetric-key and public-key encryption (and verification) of log entries are very common confidentiality techniques proposed in the literature. Unfortunately, these schemes are becoming less useful in cloud-based applications. There is a need for robust access control mechanisms that enable dynamic user addition and revocation with minimal overhead. In other words, continuously re-encrypting a subset of the log database should be avoided. Both symmetric- and public-key cryptosystems suffer in that access policies must be tied directly to keys used for encryption and decryption. If the access policy for a set of log messages needs to be changed, then both the keys used to encrypt and decrypt such log entries will need to be re-generated and distributed, and the entries must also be re-encrypted. Both of these tasks can be very expensive.

In addition, symmetric-key cryptosystems require keys to be shared among users who need access to the same set of logs. This requires a secure and comprehensive key management and distribution scheme and supporting policy. In a similar vein, public-key cryptosystems (e.g. RSA and ElGamal) suffer from the extra data transfer and storage requirements for large cryptographic keys and certificates. There may be insufficient resources to maintain a public-key infrastructure (PKI) for managing keys and digital certificates for all users.

In terms of log file integrity, aggregate signature schemes that support forward secrecy through the use of symmetric- and public-key cryptosystems are also becoming outdated

[19]. Symmetric-key schemes may promote high computational efficiency for signature generation, but they do not directly support public verifiability for administrators and auditors. This means that robust key distribution schemes or the introduction of a trusted third party (TTP) are needed to ensure that all required parties can verify the necessary log data. Such schemes also suffer from relatively high storage requirements and communication overhead. Public-key schemes have similar issues, as the increased key size leads to even larger storage requirements and less computational efficiency. Also, public-key schemes introduce the need for a trusted certificate authority to grant certificates for all parties that sign log information. One time-tested technique for supporting log file integrity is the use of authenticated hash-chains [17], which will be the focus of this paper.

Collectively, we see that a balance between encryption and signature generation and verification performance is needed to support the unique scalability and resource usage requirements for cloud-based applications. Furthermore, the selected cryptographic primitives to encrypt, sign, and verify data must not exacerbate the problem of dynamically changing access control policies and user privileges. Role-based Access Control (RBAC), which first gained popularity in the mid 1990s [15] [8] and was later proposed as a standard for the National Institute of Standards and Technology in 2001[9], is an increasingly popular access control policy that enables users to be associated with roles that change less frequently. In the context of maintaining the confidentiality of log messages generated by many users, RBAC surpasses traditional mandatory and discretionary access control (MAC and DAC) [1].

More recently, attribute-based access control (ABAC) [18] [3] [20] has been developed to provide fine-grained access control to sensitive data. It is common practice to specify user roles as attributes in this access control scheme, thus enabling the benefits of RBAC with fine-grained access control. Attribute-based encryption (ABE) [10], a new cryptographic scheme that uses user attributes (or roles, in this context) to maintain the confidentiality of user-sensitive data, has an appealing application to logging systems maintained in the cloud and is capable of satisfying the aforementioned confidentiality requirements.

In this paper we address all of the aforementioned issues with ABLS, an attribute-based logging system designed to support automated audits of encrypted audit trails (log data) based on user-defined security policies. Access to sensitive log information is enforced using ciphertext-policy attribute-based encryption (CP-ABE) with a minimal number of log-related roles, and thus a small number of attributes, to avoid the problem of increasing encryption computational complexity with attribute explosion. We present the preliminary design of ABLS and discuss how audit trails are constructed, automated audit tasks are defined and specified, and how the system may be used in practical applications.

The paper is organized in a top-down fashion, starting with the structure of log data and corresponding ability to define automated audit tasks. Using this foundation, we then introduce the relational data model used to persist log information, followed by the cryptographic access control mechanisms used to maintain the confidentiality of log data and audit trails. Finally, we conclude with a practical use case for ABLS in the context of healthcare organizations.

2. SECURE LOGGING REQUIREMENTS

The most fundamental requirements for a secure logging system is that it provides log data integrity and confidentiality. In the context of a secure logging system, integrity is based on the forward-secure stream integrity model. This model of integrity guarantee that log data cannot be forged or rearranged within the stream of log messages, and resilient against attacks that try to recover old keys after a machine has been compromised (i.e. forward secure).

Research into this problem has since revealed that a variety of other realistic requirements also exist, including a resilience to truncation and delayed detection attacks, minimal reliance on an on-line server for log storage and verification, and storage efficiency. For completeness, truncation and delayed detection attacks are defined below, based on explanations presented in [13].

- **Truncation attack** - An attack in which a set of log entries residing on an untrusted log servers is truncated, or shortened so as to remove suspicious events, without being detected by synchronization with the trusted log server.
- **Delayed detection** - An attack in which log entries on the untrusted log servers can be modified by an attacker who possesses the authentication key A_i after compromising the system between log entries L_i and L_{i+1} . With this key, the attacker can easily change entries in the log. However, once the trusted machine synchronizes with the untrusted log server to check the integrity of the log messages, this attack is immediately detected. Of course, significant damage could have already been done during this time window when the log server is compromised and it is synchronized with the trusted machine.

3. LOG FILE RELIABILITY

Log file reliability, which reduces to log integrity, is achieved through hash chains and message-authentication codes, which is a technique first introduced by Schneier et al [17] and motivated by Bellare et al [5]. Each log entry is a five-tuple element that contains the generating source information (U, S) , the encrypted payload C of the log data D , a hash digest that provides a link between the current and previous hash chain entries X , and an authentication tag for the digest Y . Formally, each log entry L_i is built using the following protocol:

$$\begin{aligned} X_i &= H(X_{i-1}, E_{PK}(D_i)) \\ Y_i &= HMAC_{A_i}(X_i) \\ L_i &= (U_i, S_i, E_{PK}(D_i), X_i, Y_i) \end{aligned}$$

In this scheme the X_i elements are used to link together consecutive entries in the hash chain. Similarly, the Y_i elements are used to provide authentication for the X_i element using an authentication tag for the X_i entries. Also, the initial value for the MAC key A_0 is randomly generated when a user session is created.

In order to prevent truncation and deletion attacks, a single entity T_i for the entire log chain is updated as the log chain is iteratively constructed. Formally, T_i is computed as follows:

$$\begin{aligned} T_i &= HMAC_{B_i}(L_i, T_{i-1}) \\ T_0 &= HMAC_{B_0}(L_i, 1) \end{aligned}$$

B_0 is a secret symmetrical key that is randomly generated when the log chain for a user session is initialized. This key is evolved with a pseudorandom function H as follows:

$$B_{i+1} = H(B_i)$$

For forward-security, this entity creation could be replaced with a publicly-verifiable Forward Secure Sequential Aggregate (FssAgg) signature scheme backed by a trusted certificate authority to support forward-secure stream integrity. We refer the reader to [12] for a more thorough treatment of FssAgg signature schemes and their role in updating the single log chain entity.

3.1 Verification Modes

The ABLS log construction scheme enables two different modes of verification to be implemented, each of which has different integrity and performance guarantees. The first mode of verification requires one to walk the log chain, computing the digest X_i and comparing it to the value stored in the database. However, this method does not guarantee the integrity of the log chain. The second mode of verification requires a trusted verifier task \mathcal{V} to use the initial hash chain key A_0 and entity key B_0 to walk the log entries, computing both X_i and Y_i , and comparing them against the values stored in the database. While this mode does not lend itself to public verifiability, it guarantees the integrity of the log chain if a forward-secure MAC is used.

4. LOG ACCESS CONTROL

Access control for log data is enforced using ciphertext policy attribute-based encryption (CP-ABE), a new encryption scheme that supports complex access policies that specify which secret keys can be used for decryption [6]. In CP-ABE, secret keys are analogous to sets of attributes, and access policies are defined using tree-like access structures of logical AND and OR gates, where each leaf in the tree is an attribute. Implementations of CP-ABE schemes are usually based on the construction of a bilinear mapping between two elliptic curve groups [6] [11]. We define both of these terms in the following sections.

4.1 Mathematical Foundations

Definition Let \mathbb{F}_p be a finite field where $p > 3$ is a prime, and $a, b \in \mathbb{F}_p$ such that

$$4a^3 + 27b^2 \neq 0 \pmod{p} \quad p \in \mathbb{F}_p$$

An elliptic curve $E[\mathbb{F}_p]$ is the set of solutions (x, y) to the equation

$$y^2 = x^3 + ax + b \pmod{p} \quad p \in \mathbb{F}_p[x],$$

together with the point at infinite 0.

Definition Let G_1 and G_2 be cyclic groups of prime order p and g a generator of G_1 . We say that e is a *bilinear map* defined as $e : G_1 \times G_2$, where $|G_1| = |G_2| = p$. This bilinear map satisfies the following properties:

- **Bilinearity:** For all $u, v \in G_0$ and $a, b \in \mathbb{Z}_p$, we have $e(u^a, v^b) = e(u, v)^{ab}$
- **Non-degeneracy:** $e(g, g) \neq 1$
- **Computability:** Both G_1 and G_2 are efficiently computable

4.2 Ciphertext Policy Attribute-Based Encryption

In the original construction of the CP-ABE scheme, Bethencourt et al. [6] defined five different procedures used in the cryptosystem: *Setup*, *Encrypt*, *KeyGeneration*, *Derypt*, and *Delegate*. We define each of these procedures as follows:

- **Setup** - This procedure takes the implicit security parameter as input and outputs the public and master keys PK and MK .
- **Encrypt**(PK, M, \mathbb{A}) - This procedure will encrypt M , a plaintext message, to produce a ciphertext CT such that only a user that possesses a set of attributes that satisfies the access structure \mathbb{A} will be able to decrypt the message. The encryption process embeds \mathbb{A} into the ciphertext.
- **KeyGeneration**(MK, \mathbb{S}) - This procedure generates a private key SK using the master key MK and set of attributes S that describe the private key.
- **Decrypt**(PK, CT, SK) - This procedure decrypts the ciphertext CT using the provided secret key SK to return the original message M . Decryption is only successful if the set S of attributes, which is associated with the key SK , satisfies the access policy embedded within the ciphertext (which is part of the access structure \mathbb{A}).
- **Delegate**(SK, \tilde{S}) - This procedure outputs a secret key \tilde{SK} for the set of attributes \tilde{S} , where $\tilde{S} \subset S$, the set of attributes associated with the secret key SK .

4.3 Access Policy Definition

A major component of ABLS is the policy engine, which maps access policies defined on an event basis to the corresponding access tree used for encryption. For example, an access policy might state that only User XYZ, or physician assistants or nurse practitioners from Medical Group A, are allowed to access data associated with an event E . The corresponding access tree for this policy is shown in Figure 1.




Figure 1: Access tree for a policy that only enables access to user XYZ, Medical Group A physician assistants, and Medical Group A nurse practitioners.

Our logging scheme makes the assumption that events, and the access policies for all data associated with such events, are well defined, which is often the case with organizations that must comply with federal legislation like HIPAA [4]. With this assumption, the behavior of the policy engine in

generating access policies is dependent on administrator-defined policy rules for events of interest and the corresponding attributes of system users and data that trigger such events. In this way, policy rules are coupled to events so that policies for access are generated based on the type of event that occurred and the user who is requesting access to such information.

4.4 Key Generation and Management

Pairing-based cryptography is computationally expensive, and under the assumption that ABLS might be subject to very heavy traffic loads at any particular time, the overhead of encrypting data to be stored in the database should be as minimal as possible. Therefore, each unique policy that is needed to encrypt a log message is associated with a symmetric key, which is in turn encrypted using CP-ABE and then serialized to be stored in the key database. This design enhancement enables increased throughput without sacrificing the level of confidentiality granularity that is needed for each log entry. However, should an unencrypted policy key for a given user's session become compromised, the remaining entries in that log database are at risk of being compromised.

The basic procedure for encrypting a log entry is shown in Algorithm 1. Once encrypted, the ciphertext is stored in the database with the rest of the information necessary to continue the log chain for a given user's session.

Algorithm 1 Log entry encryption

Require: An unencrypted log entry L_i for session S_j of user U_k

- 1: Let P be the access control policy for the message of L_i , as determined by the `PolicyEngine`
- 2: **if** The symmetric key K for (U_k, S_j) has not been generated for P **then**
- 3: Generate K and encrypt it with the CP-ABE encryption module using P , yielding K_E
- 4: Persist K_E to the key database
- 5: **else**
- 6: Query the database for K_E , the encrypted key for policy P .
- 7: Decrypt K_E using the attributes of user U_k , yielding K
- 8: Encrypt L_i with AES-256 using K , yielding $E(L_i, K)$
- 9: Persist $E(L_i, K)$ to the log database

In order to improve the performance of the logger, the per-policy symmetric keys used for encryption for a user session are kept in memory until the session has been closed. This avoids the need for the logger to query the database for the key when a new log message arrives.

By default, the secret keys used for decryption are never cached in the system's local memory. Since it is expected that log entries will be read much less frequently than they will be written, such keys are generated on demand by querying the appropriate database. Furthermore, the key generation process can be done in two ways. For policy rules that limit the access to only the generating user, only a single query to the attribute database is required to establish the user's secret key and then decrypt the data for all log entries corresponding to that rule. With this key the user may decrypt these entries offline without the need to query the policy engine for the appropriate access rights.

Conversely, for access policy rules that embed attributes for colleagues or other users related to the source user, the policy engine must first query the user database to ensure the requesting user meets the relationship criteria set in place by the policy. Then, if this is successful, the policy engine will grant the appropriate secret key to the requesting user. The tradeoff is that, while an online TTP is needed for such colleagues to access the log entry contents, it is significantly easier for the system administrators to manage who has access to specific log entries aside from the original source user. Simply modifying the user's relationships in the system database is sufficient to revoke access from certain colleagues.

In order to maintain the security of the system at runtime, it is necessary to cycle the master and public keys associated with the encryption scheme. Our current system does not support this feature, but there are two ways that it could be implemented. The first way is to persist the old master and public keys to a safe location that could be called during auditing and verification if needed. The second way is to re-encrypt the entire log database with the new master and public key. Unfortunately, this would not only require the system to be brought offline during the update (in order to avoid synchronization issues with live traffic), but it would also mean that the new master and public key serve as a single point of failure for the entire database if compromised. Therefore, future releases of this system plan to implement the first approach to manage keys. It would be best to determine the key cycle lifetime based on empirical data associated with the growth of the log database. Intuitively, in order to maintain auditing and verification efficiency, the cycle frequency should be defined as a monotonically increasing function that is proportional to the growth of the database.

5. STRUCTURED LOG DATA FOR AUTOMATED AUDITS

Current technological solutions for automated audits in organizations where the usage of audit trails is required for compliance activities are considered to be ineffective [7]. In order to support automated audits of audit trails, it is necessary to specify a formal structure for this data that matches organizational security policies. Therefore, a major motivating factor for the ABLS log data structure comes from realistic security policies. In this context, we make the assumption that a security policy can be stated as a set of *negative* requirements. For example, one such requirement might be that a doctor is not allowed to change their patient's address. In order to conduct an automated audit for violations of this policy, we first translate this semantically-rich requirement into a language whose structure can be easily mapped to a relational data model. This enables us to leverage the power of structured query languages (i.e. SQL) to search for policy violations.

One solution for parsing security policy requirements into relational data is to define a grammar for producing requirement strings from a set of non-terminals that correspond to relations. Using the NIST RBAC model of access control as motivation [9], we specify this set of non-terminals and relations to be the set identifiers USER, OBS, OPS, and AFFECT-EDUSERS. These finite sets are minimal enough to allow the specification of most security policies, thus making it suitable for our needs.

LAudit, a simple context-free grammar that is built on these

relations, is shown below.

```

LAudit ::= USER OPS
        | USER OPS OBS
        | USER OPS OBS USER

```

In this context, USER, OPS, and OBS are all finite sets composed of the users, operations, and objects of a system, as specified by the NIST RBAC model [16]. While simple, this language effectively captures the “who” and “what” of log events. ABLS is capable of appending a timestamp to every that it receives, which rounds off the log event with “when” information.

ABLS clients must submit log messages according to a pre-defined schema that captures all of the information in LAudit. A JSON schema that can be used for constructing log messages is shown below.

```

[
  {
    user : int ,
    session : int ,
    action : int (or String),
    object : int (or String),
    affectedUsers : [int]
  }
]

```

6. RELATIONAL MODEL AND PRIVACY IMPLICATIONS

In order to capture the audit information in a relational model to enable efficient and automated queries, the events, actions, objects, and affected users are all coupled to the incoming log data. The resulting relational model is shown in Figure 2.

This schema only corresponds to the log database in ABLS. There are in fact four databases altogether that must be maintained by ABLS: the log, key, user, and policy database. The log database maintains all information in the log chain for every single user and session pair. The key database stores the cryptographic keys that were used to construct such log chains. The user, and policy databases store user information and policy rules for ABLS, respectively.

In order to link the entries in the log tables to their corresponding verification and encryption keys in the key database, common user and session IDs are used (though not as the primary key for the tables since they do not satisfy the uniqueness property). However, storing user and session information in plaintext may lead to a privacy violation if the database is compromised. Therefore, using a technique similar to the “onion encryption” design in CryptDB [14], this information is now deterministically encrypted before being stored in the database.

This procedure works by encrypting the user and session attributes with a symmetric key generated from the logger’s master key salted by the target table identifier. In mathematical terms, the encrypted user and session IDs, $[U_i]$ and $[S_j]$, stored in table T are generated as follows.

$$\begin{aligned}
[U_i] &= E(M_k || H(T), U_i) \\
[S_j] &= E(M_k || H(T), S_j)
\end{aligned}$$

In this context, M_k is the master key for the logger. Using the table identifier as a salt to the master key enable ensures that tables do not share any common information about the user, which helps prevent against inference attacks in the event that the database servers are compromised. Furthermore, this enables verifiers, who will have access to M_k through the key manager, to decrypt log entries and recover the user identifier so that they may check the contents of the other databases as needed.

In the log data model, all Action and Object records are stored in plaintext. These tables store elements of a finite set, and encrypting them would not derail a determined attacker. However, all information about affected users is encrypted (masked) using the aforementioned technique. As such, an attacker can infer information about what types of objects were operated upon, but they cannot determine the specifics of these actions or the users who performed them without compromising the ABLS master key M_k . We feel as though this strikes a good balance between robust audit specification, reasonable measures of audit and log efficiency, user privacy, and log security.

7. DEPLOYMENT

ABLS is designed to be a centralized logging system backed by a set of distributed databases. A context diagram for the ABLS deployment scheme is shown in Figure 3.

Based on the purpose of each piece of data used in the log, it is best to physically separate databases that store data of different security classes rather than rely on a single, segregated database that uses MAC with polyinstantiation to protect data of different security classes. Of course, access control and authentication mechanisms for all of the database servers is to be enforced at the operating system level, thus prohibiting immediate access to all unauthorized users other than the internal tasks (i.e. logger, verifier, policy engine, etc) within an ABLS instance.

8. ANALYSIS

The prototype ABLS system was written in Python, utilizing the Charm cryptography package [2] for pairing-based cryptographic primitives. A major concern for this new architecture was its scalable performance while processing large amounts of log data. Therefore, we conducted experiments to measure the encryption overhead from events that have uniform access policies, as well as the encryption overhead from events that have different access policies. These two experiments were conducted to gauge the overhead of that results from log encryption and querying the policy engine for generating new policies. The results for these two experiments are shown in Figures 4 and 5.

To measure the storage overhead for the log data we compared the physical size of an empty SQLite database against one that contained log data for three different policies, each with different sets of attributes. The results of this comparison are shown in Table ?? . Based on this analysis, we concluded that each log entry contributed approximately 1.2Kb of information to the log database, which is largely due to the highly conservative data types in the relational model. This overhead could be drastically reduced if the model was adjusted to match the cryptographic primitives used in the logging scheme and avoiding the use of BLOB attribute types.

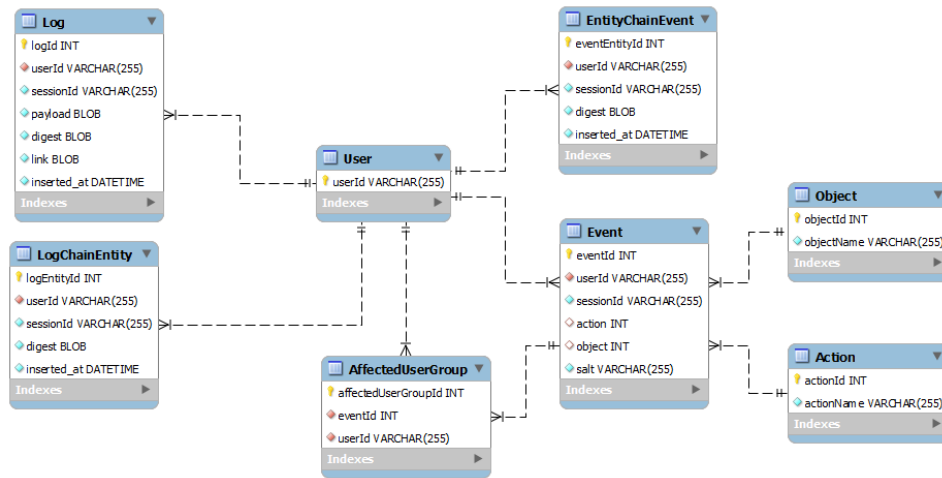


Figure 2: ER diagram depicting the relational schema for the log data model.

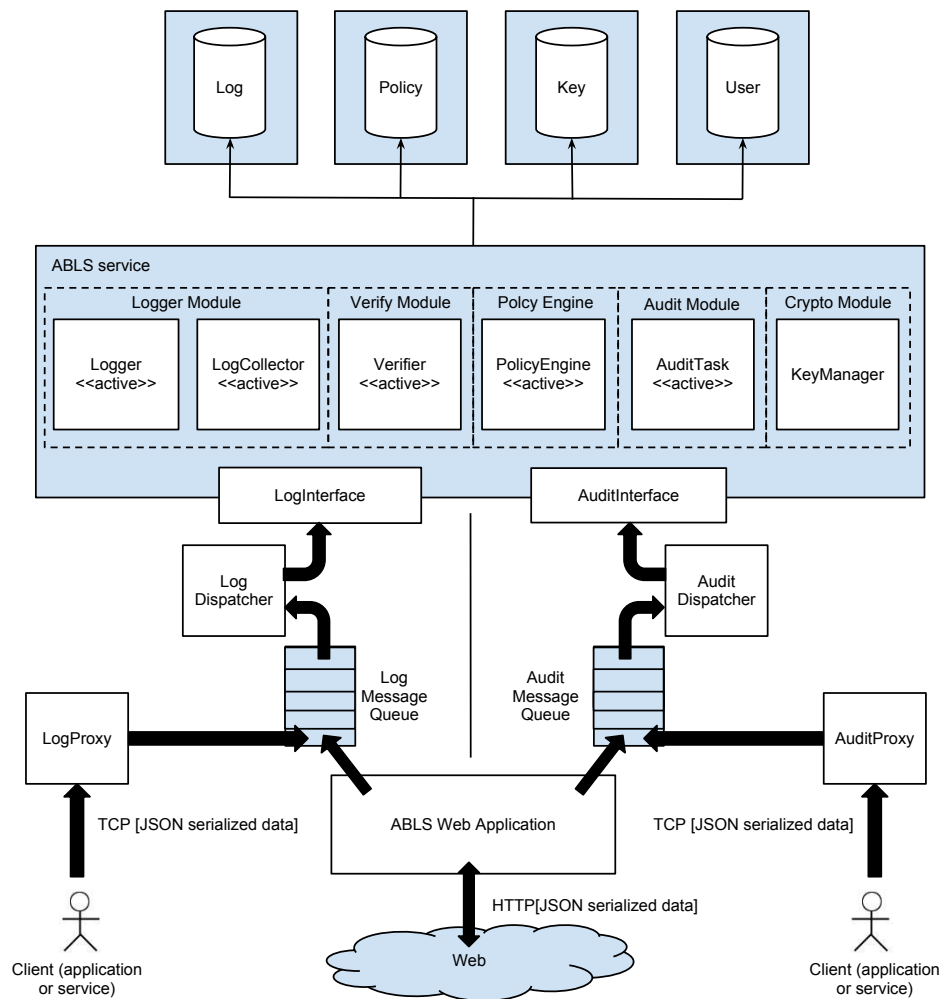


Figure 3: A high-level deployment diagram for ABLs.

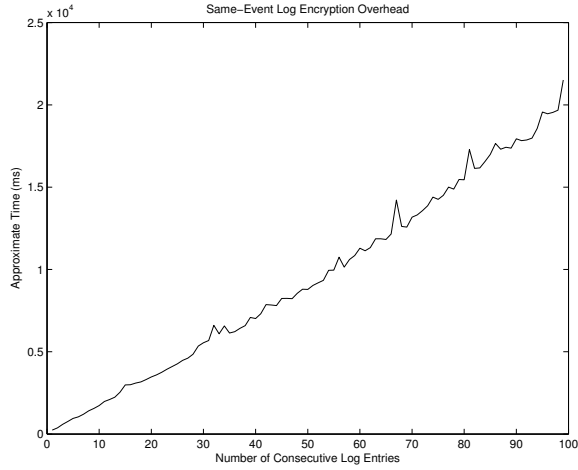


Figure 4: Log encryption performance for data in a single user session that all have the same access policy.

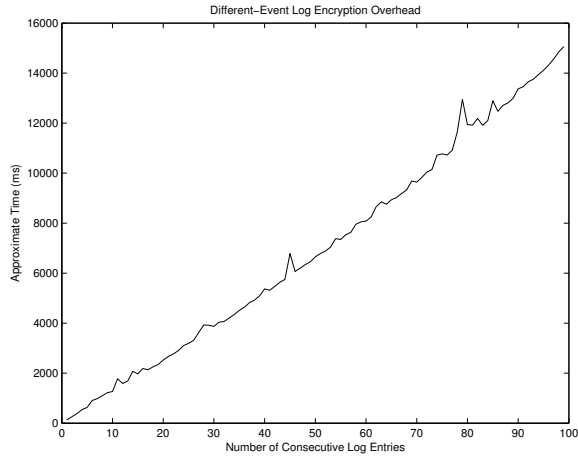


Figure 5: Log encryption performance for data in a single user session that all have different same access policies.

9. APPLICATIONS

ABLS would be particularly useful in the healthcare domain, where organizations are required to comply with HIPAA policies and regulations. Non-repudiation is a critical element that is needed to support these compliance efforts. Therefore, configuring ABLs for use in these organizations can help securely collect and persist relevant log information, and also automate the auditing process to reinforce the manual efforts done by security administrators. In such an environment, healthcare web applications would be configured to send all relevant user events to an ABLs instance in the cloud via the HTTP API. All clients could be authenticated using a standard password-based scheme that is connected the organization’s user database.

Similarly, all dedicated health machines and terminals in hospitals and physician offices would maintain continuous data streams to an ABLs instance and send all log data over a TCP socket using an API similar to that which is exposed to web applications. These dedicated entities could be au-

Table 1: Storage overhead for storing the log data

| Log Database Contents | Storage Size (Kb) |
|--------------------------------------|-------------------|
| None | 12 |
| 100 log entries - minimal attributes | 134 |
| 100 log entries - medium attributes | 147 |
| 100 log entries - maximum attributes | 149 |

thenticated using two-way SSL authentication so as to allow authentication to take place using automatically exchanged SSL certificates.

Security administrators could review the organization’s required security policies and encode them using the popular eXtensible Access Control Markup Language (XACML), which could be used to configure the ABLs policy engine for encrypting all incoming log data. Similarly, the security administrators could define automated audit tasks that ensure policy violations do not occur. The security administrator would also be responsible for configuring the frequency at which automated audit and verification tasks and run. Although the current ABLs prototype does not currently export policy violations to a dashboard of any kind, in the event that this feature was added, the security administrator would also be responsible for configuring the connection between ABLs and the dashboard. Finally, depending on the rate at which the ABLs databases grow, the security administrator should also specify the parameters for archiving all audit data in a remote location.

10. CONCLUSION

In this paper we presented ABLs, a secure logging system for the cloud that supports automated audits and dynamic, attribute-based log file encryption. We discussed the log integrity and confidentiality properties supported by ABLs, and how they map to the underlying system architecture and relational data model. We then presented a sample deployment scheme for an ABLs instance in the cloud, which included the client interfaces for submitting log data and executing audit queries. We finished with a performance analysis of the current ABLs prototype and discussion of employing ABLs in the healthcare domain. This paper revealed that achieving automated audits is a non-trivial challenge that must strike a balance between system performance and storage overhead. While the current ABLs design provides efficient automated audits to help support non-repudiation, there are several fundamental design changes (i.e. the application of FssAgg signature generation algorithms to encrypted log chains) that can be explored for further improvement.

11. REFERENCES

- [1] M. Abrams, K. Eggers, L. LaPadula, and I. Olson. A generalized framework for access control: An informal description. In *Proceedings of the 13th National Computer Security Conference*, pages 135–143, 1990.
- [2] J. Akinyele, M. Green, and A. Rubin. Charm: A framework for rapidly prototyping cryptosystems.
- [3] H. Alipour, M. Sabbari, and E. Nazemi. A policy based access control model for web services. In *Internet Technology and Secured Transactions (ICITST)*, 2011 *International Conference for*, pages 472–477. IEEE, 2011.

- [4] G. Annas. Hipaa regulations—A new era of medical-record privacy? *New England Journal of Medicine*, 348(15):1486–1490, 2003.
- [5] M. Bellare and B. S. Yee. Forward integrity for secure audit logs. Technical report, 1997.
- [6] J. Bethencourt, A. Sahai, and B. Waters. Ciphertext-policy attribute-based encryption. In *Proceedings of the 2007 IEEE Symposium on Security and Privacy*, SP '07, pages 321–334, Washington, DC, USA, 2007. IEEE Computer Society.
- [7] P. I. T. A. Committee et al. Revolutionizing health care through information technology. Arlington, VA: *National Coordination Office for Information Technology Research and Development*, 2004.
- [8] F. David and K. Richard. Role-based access controls. In *Proceedings of 15th NIST-NCSC National Computer Security Conference*, volume 563. Baltimore, Maryland: NIST-NCSC, 1992.
- [9] D. Ferraiolo, R. Sandhu, S. Gavrila, D. Kuhn, and R. Chandramouli. Proposed nist standard for role-based access control. *ACM Transactions on Information and System Security (TISSEC)*, 4(3):224–274, 2001.
- [10] V. Goyal, O. Pandey, A. Sahai, and B. Waters. Attribute-based encryption for fine-grained access control of encrypted data. In *Proceedings of the 13th ACM conference on Computer and communications security*, pages 89–98. ACM, 2006.
- [11] P. Junod and A. Karlov. An efficient public-key attribute-based broadcast encryption scheme allowing arbitrary access policies. In *Proceedings of the tenth annual ACM workshop on Digital rights management*, DRM '10, pages 13–24, New York, NY, USA, 2010. ACM.
- [12] D. Ma. Practical forward secure sequential aggregate signatures. In *Proceedings of the 2008 ACM symposium on Information, computer and communications security*, ASIACCS '08, pages 341–352, New York, NY, USA, 2008. ACM.
- [13] D. Ma. Practical forward secure sequential aggregate signatures. In *Proceedings of the 2008 ACM symposium on Information, computer and communications security (ASIACCS'08)*, pages 341–352, 2008.
- [14] R. A. Popa, C. M. S. Redfield, N. Zeldovich, and H. Balakrishnan. Cryptdb: processing queries on an encrypted database. *Commun. ACM*, 55(9):103–111, Sept. 2012.
- [15] R. Sandhu, E. Coyne, H. Feinstein, and C. Youman. Role-based access control models. *Computer*, 29(2):38–47, 1996.
- [16] R. Sandhu, D. Ferraiolo, and R. Kuhn. The nist model for role-based access control: towards a unified standard. In *Symposium on Access Control Models and Technologies: Proceedings of the fifth ACM workshop on Role-based access control*, volume 26, pages 47–63, 2000.
- [17] B. Schneier and J. Kelsey. Secure audit logs to support computer forensics. *ACM Transactions on Information and System Security (TISSEC)*, 2(2):159–176, 1999.
- [18] H. Shen and F. Hong. An attribute-based access control model for web services. In *Parallel and Distributed Computing, Applications and Technologies*, 2006. PDCAT'06. *Seventh International Conference on*, pages 74–79. IEEE, 2006.
- [19] A. A. Yavuz and P. Ning. Baf: An efficient publicly verifiable secure audit logging scheme for distributed systems. In *Proceedings of the 2009 Annual Computer Security Applications Conference*, ACSAC '09, pages 219–228, Washington, DC, USA, 2009. IEEE Computer Society.
- [20] J. Zhu and W. Smari. Attribute based access control and security for collaboration environments. In *Aerospace and Electronics Conference*, 2008. NAECON 2008. *IEEE National*, pages 31–35. IEEE, 2008.