ABLS - An Attribute Based Logging System for the Cloud

Christopher A. Wood
Department of Computer Science
`caw4567@rit.edu`

ABSTRACT

User-based non-repudiation is an increasingly important property of cloud-based applications. It provides irrefutable evidence that ties system behavior to specific users, which in turn enables the strict enforcement of organizational security policies. System logs, which can be used to construct audit trails, are typically used as the basis for this property. Thus, the effectiveness of system audits based on log files reduces to the problem of maintaining the integrity and confidentiality of log files. We present the design and implementation of an attribute-based logging system, ABLS, that is capable of building secure log files. In doing so, we also present some of the benefits of ciphertext-policy attribute-based encryption (CP-ABE) to solve a variety of log design issues. In addition, we also present the design of an automated auditing procedure that is dynamically configurable by its users to aid in security policy enforcement.

In this paper we discuss ABLS, an attribute-based logging system that supports ciphertext-policy attribute-based encryption (CP-ABE) [5] and authenticated hash-chain constructions for log file confidentiality and integrity, respectively. ABLS was designed for the cloud and is capable of addressing the aforementioned problems with secure logging. We first start with progress that has been made since the first phase, which includes versions 1 and 2 of ABLS, and then present a plan of attack for the next phase in the project.

1. ABLS-V1 COMPLETED WORK

For Phase 2 of the project, I focused on finishing remaining work from the first phase (including an instance-wide key manager to help verifiers work with the log information and finishing the database masking procedure), implementing the log collector to collate log messages to be inserted into the database, and finishing the log query capability. The specifics of each of these tasks is outlined in the following sections. The design and implementation of the automated audit tasks proved to be more difficult than originally anticipated and required a complete re-design of the relational model and structure of log messages. As such, this feature was omitted to ABLS-V2, and is described later in this document.

1.1 Deployment

ABLS is designed to be a centralized logging system backed by a set of distributed databases. A context diagram for the ABLS deployment scheme is shown in Figure 1.

Based on the purpose of each piece of data used in the log, it is best to physically separate databases that store data of different security classes rather than rely on a single, segregated database that uses MAC with polyinstantiation to protect data of different security classes. Of course, access control and authentication mechanisms for all of the database servers is enforced at the operating system level, thus prohibiting immediate access to all unauthorized users other than the internal tasks (i.e. logger, verifier, policy engine, etc) within an ABLS instance.

In the current ABLS prototype, all database servers are separated as individual SQLite database files. Once the deployment platform is selected and properly configured, these will be replaced with instances of MySQL databases running on separate servers.

1.2 Hybrid Key Management

Pairing-based cryptography is computationally expensive, and under the assumption that ABLS might be subject to very heavy traffic loads at any particular time, the overhead of encrypting data to be stored in the database should be as minimal as possible. Therefore, each unique policy that is needed to encrypt a log message is associated with an AES-256 symmetric key, which is in turn encrypted using CP-ABE and then serialized to be stored in the key database. This design enhancement enables increased throughput without sacrificing the level of confidentiality granularity that is needed for each log entry. However, should an unencrypted policy key for a given user's session become compromised, the remaining entries in that log database are at risk of being compromised. The Charm crypto package allows all cryptographic objects (which tend to be nested Python dictionaries and other complicated data structures) to be serialized to byte representations for database persistence.

The basic procedure for encrypting a log entry is shown in Algorithm 1. Once encrypted, the ciphertext is stored in the database with the rest of the information necessary to continue the log chain for a given user's session.

In order to improve the performance of the logger, the per-policy symmetric keys for a user session are kept in memory until the session has been closed. This avoids the need for the logger to query the database for the key when a new piece of data to be inserted into the log.

Also, in order to ensure that every encryption module (cipher) contained within loggers, verifiers, and database shims uses the same master key M_k, a key manager singleton object was implemented and shared among all ABLS entities that require the master

Figure 1: A high-level depiction of the deployment for an ABLS instance, where each box represents a unique runtime environment (i.e. a unique server).

Algorithm 1 Log entry encryption

Require: An unencrypted log entry L_i for session S_j of user U_k
1: Let P be the access control policy for the message of L_i, as determined by the `PolicyEngine`
2: **if** The symmetric key K for (U_k, S_j) has not been generated for P **then**
3:     Generate K and encrypt it with the CP-ABE encryption module using P, yielding K_E
4:     Persist K_E to the key database
5: **else**
6:     Query the database for K_E, the encrypted key for policy P.
7:     Decrypt K_E using the attributes of user U_k, yielding K
8: Encrypt L_i with AES-256 using K, yielding $E(L_i, K)$
9: Persist $E(L_i, K)$ to the log database

---

key. Upon creation, an encryption module will register itself with the key manager so that they receive any key changes in the event that the master (or public) key changes.

1.3 Database Design

As shown in section 1.1, there are five main databases that must be maintained by ABLS: the log, key, user, audit_user, and policy database. The log database maintains all information in the log chain for every single user and session pair. The key database stores the cryptographic keys that were used to construct such log chains. The user, audit_user, and policy databases store user information and policy rules for ABLS, respectively. In the current prototype of ABLS, the policy database is not used internally. Instead, all event rules are hard-coded into the `PolicyEngine`.

In order to link the entries in the log tables to their corresponding verification and encryption keys in the key database, common user and session IDs are used (though not as the primary key for the tables since they do not satisfy the uniqueness property). However, because such user and session information is stored in plaintext, their existence may lead to a violation of user privacy. Therefore, using a technique similar to the "onion encryption" design in CryptDB [6], this information is now deterministically masked before being stored in the database.

This masking procedure works by encrypting the user and session attributes with a symmetric key generated by the logger's master key salted by the target table identifier. In mathematical terms, the encrypted user and session IDs, $[U_i]$ and $[S_j]$, stored in table T are generated as follows.

$$[U_i] = E(M_k \| H(T), U_i)$$
$$[S_j] = E(M_k \| H(T), S_j)$$

In this context, M_k is the master key for the logger. Using the table identifier as a salt to the master key enable ensures that tables do not share any common information about the user, which helps prevent against inference attacks in the event that the database servers are compromised. Furthermore, this enables verifiers, who will have access to M_k through the key manager, to decrypt log entries and recover the user identifier so that they may check the contents of the other databases as needed.

The relational database design was modified in order to support the new symmetric key management scheme and provide enhanced security for the stored data. In particular, rather than storing the cryptographic keys in the same database as the log information, these two databases are now segregated and are expected to be deployed on different database servers. By hardening both databases, a malicious attacker would now need to circumvent the access control mechanisms protecting two physically disjoint databases, rather than just one, in order to recover log decrypted messages. In terms of the actual implementation, this change caused many cascading errors throughout the code base which needed to be fixed before development continued.

Also, to better support audits that use the log database, a timestamp field was added to each table as a required attribute. Not only does such information capture the exact timing of critical system events, it acts as a protection mechanism in the event that log messages are inserted into the database out of order. Also, it is important to note that each and every timestamp for a log message is generated when a database entry is generated to be sent to the log collector.

1.4 Log Querying

In order to make ABLS fulfill its purpose as a secure logging system, log query support will be added to the audit module. Such

query support will be exposed to clients through a lightweight API that is used via the audit protocol. To start, only select operations parameterized by the user ID or both the user and session IDs will be supported. As an example, the API might have the following function signatures.

```
selectByUser(int uid)
selectByUserAndSession(int uid, int sid)
```

Clients will connect to the audit proxy, which is similar to the log traffic proxy, and issue commands according to the audit protocol to invoke either one of these functions. Client authentication is expected to be completed at a later point in time as outlined in Section 3.1. Assuming the client has been authenticated, they can issue commands wrapped in JSON strings with the following format to the audit proxy.

```
{
    "command" : command_id
    "params" : csv_list
}
```

Upon receiving and parsing an incoming command, the client handler, which is spawned to handle each client connection, will issue the appropriate command to the log module to select the appropriate log message contents from the database. However, since each client is expected to be a user for the ABLS instance, their attributes will be used to filter the resulting database rows that are returned. Thus, the audit proxy and client handlers will act as a reference monitor that enforces access control to the log database through the log query protocol and API.

Unlike the log proxy, it is expected that the clients for the audit proxy will be authenticated using passwords. Thus, as part of supporting client authentication, a proper password storage scheme was implemented in which the hash of password digests salted by a random string are persisted into the database. In this way, the user password is never sent in plaintext over the network. Also, the salt is stored in plaintext in the same audit_user record. This is a common design decision for web applications.

1.5 Log Collection

As part of the proposed design, all log messages generated by the logger to be inserted into the database will be handed off to a common log collector to do so. This log collector is implemented as a separate thread running within the ABLS process and will assume a portion of the database communication overhead from the logger to free up cycles for processing more log messages. Each logger instance points to the same log collector thread. Furthermore, all log collector tasks are spawned when the ABLS instance first loads.

1.6 Bootstrapping, Test Enhancement, and Deployment

In order to streamline the test and deployment phases of development, a bootstrap script was implemented to configure new (empty) versions of the local SQLite databases. The main executable script (Main.py) was also modified to support development and production modes of operation, in which the development mode clears the contents of every database and then proceeds to insert false user data into the users table to begin the logging process. The test driver program, LogProxyDriver.py, was then modified accordingly to use the default data contained within the database. With these changes, the typical process to start and interact with an ABLS instance is as follows.

1. Run the bootstrap script to create new versions of the local SQLite databases. In the actual deployment of an ABLS instance, this script would connect to the remote databases and specify their schemas accordingly. Thus, it is meant only for development purposes and should not be used on a live ABLS instance.

2. Run the main executable file (Main.py) with the -l (configure) to flag to enable development mode, in which the databases are wiped clean and replaced with a small set of fake data. If one wishes to start an ABLS instance at this point in time, they may also add the -s (start) flag to the script so that it starts the ABLS instance.

3. Run the test driver program (LogProxyDriver.py) and point it to the host and port at which the traffic proxy within the ABLS instance is listening. By default, these are "localhost" and port 9998, but they can easily be changed to any other values in the ABLS configuration file.

Aside from the initialization code, the test driver program now includes a more robust suite of tests to simulate varying traffic loads. The user interface of this program was also improved so as to aid the developers in interacting with the ABLS prototype at runtime. Given the difficulty of testing this distributed system at runtime, creating a more sophisticated test driver was crucial to the development process that enabled smoke tests to be run with minimal effort.

However, despite the increased complexity of the test driver, it does not, nor will it ever, support the ability to acquire diagnostic information from the ABLS runtime. This information is logged by the ABLS instances to the appropriate log file, and the administrators for the ABLS system can check this information at their discretion.

1.7 Log Traffic Encryption

The ABLS prototype now supports SSL encryption of all incoming log traffic to ensure the confidentiality of sensitive information as it is sent from the client to the server running the traffic proxy. Unfortunately, this is only one-way SSL authentication in which the client verifies the server. I am currently experimenting with ways to implement two-way SSL authentication so the client can be verified as well, which is the ultimate goal.

1.8 ABLS Configuration File

In order to improve the quality of the ABLS prototype, all of the hard-coded configuration strings that set up databases, specify cryptographic certificate locations, and log file output directories will be removed and placed in a configuration file to be managed by system administrators. Access to this file will be restricted to system administrators and enforced by the operating system. An example of how the configuration file might look is shown below.

```
# Network configuration paramters
abls_host = localhost
abls_logger_port = 9998
abls_audit_port = 9999

# Database configuration string
location.db.log = ~/log.db
location.db.key = ~/key.db
location.db.users = ~/users.db
location.db.audit_users = ~/audit_users.db
location.db.policy = ~/policy.db
```

1.9 Research and Literature Survey

An initial action item for this work was to investigate the possibility of replacing the relational database model with a NoSQL model. I spent a significant amount of time during week 4 investigating the benefits of making a transition to such a DBMS. In particular, I examined MongoDB, a common document-based NoSQL database. Unfortunately, my research revealed that MongoDB is not an ACID system, and thus transactions are not guaranteed to keep the database in a consistent state. In the context of ABLS, a failed transaction could put the system in a bad state. Therefore, I made the decision to keep the relational database model but hide the user-sensitive information using the techniques discussed in Section 1.3.

A comprehensive literature survey on related logging and non-repudiation work was also completed during this phase of the project. The corresponding articles that were read are included in this phase's submission package. The digest of this work is expected to be a part of the final paper.

2. ABLS-V2 DESIGN

The final proposed work item for this phase of the project was the automated audits of user information according to audit rules. Unfortunately, given the current log scheme, all context information about a log event is encrypted and stored in the log message. Even in the case that an audit task was able to decrypt a log payload and programmatically determine the details of that particular event, the audit task would need to search over every single entry in the log database when performing checks against its audit rule. In this context, the current relational data model does not support automated audits. Therefore, I chose to revisit the log generation scheme and corresponding relational model in an attempt to facilitate more efficient audits. The first step in this process was to define exactly how audit rules might be specified. To this end, I defined a custom audit language LAudit that enables users to specify audit rules for automated audit tasks. A grammar for the LAudit language is shown below.

```
LRule  ::=  USER OPS
        |   USER OPS OBS
        |   USER OPS OBS USER
```

In this context, USER, OPS, and OBS are all finite sets composed of the users, operations, and objects of a system, as specified by the NIST RBAC model [7]. While simple, this language effectively captures the "who" and "what" of log events. ABLS is capable of appending a timestamp to every that it receives, which rounds off the log event with "when" information.

With this language, ABLS clients must submit log messages according to a pre-defined schema that captures all of the information in LAudit. The JSON schema used for constructing log messages to be inserted into the log is shown below.

```
[
    {
        user : int ,
        session : int ,
        action : int (or String),
        object : int (or String),
        affectedUsers : [int]
    }
]
```

In order to capture this information in a relational model to enable efficient queries, a new Event table was added to the database schema and is backed by an EventEntityChain table for storing the digest of Event chains as they are constructed. There is a one-to-one correspondence between Event and Log records, and the security of such Event and Log information is maintained using the same hash chain construction techniques as in the preliminary design. However, for simplicity, the notion of hash chain epochs was removed, and now the log chains are strictly linear sequences of log messages that are received from clients. Also, Action, Object, and AffectedUserGroup tables were added to the database schema to store relevant information about log events as they are received by the log proxy. A high-level depiction of this new relational model is shown in Figure 2.
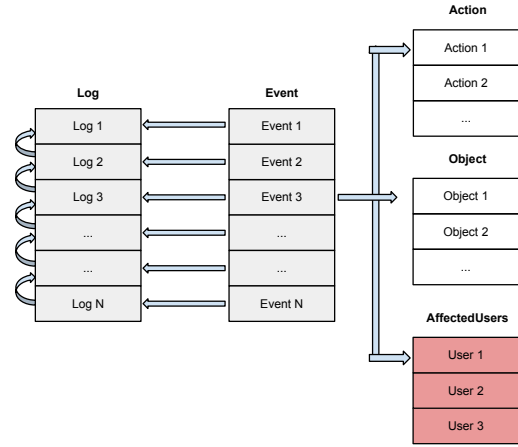

Figure 2: A high-level depiction of the new relational data model that supports automated audit tasks.

In this model, all Action and Object records are stored in plaintext. These tables store elements of a finite set, and encrypting them would not deter a determined attacker. However, all information about affected users is encrypted (masked) using the same technique discussed in Section 1.3. As such, an attacker can infer information about what types of objects were operated upon, but they cannot determine the specifics of these actions or the users who performed them without compromising the ABLS master key M_k.

Audit tasks enforce audit rules using a blacklist approach. That is, audit rules are specified using the aforementioned log message schema and then assigned to audit tasks that periodically run to see if the rules are being properly enforced. For example, an audit rule might be configured as follows:

```
[
    {
        role : "SuperUser",
        action : "Modified",
        object : "Object-X",
        affectedUsers : [1, 2, 3]
    }
]
```

Notice that the user attribute was replaced with a more abstract role attribute. Since it is rare that audit policies will be specified at the level of individual users, this enables ABLS to enforce more generic audit rules across the entire database. If an audit task was given this rule definition, it would query the Event table for all

"Modified" actions performed on objects of type "Object-X" and join the results with the events in the AffectedUserGroup table. With this data, it would check to see if any of the resulting records have affected users 1, 2, or 3 in them, and if so, generate an error or alarm.

The current implementation for ABLS-V2 supports the new log and event chaining, log message parsing, and audit task configuration using the aforementioned techniques. However, the code is mostly a proof-of-concept implementation, and is not integrated into the entire ABLS system architecture. That is, test-code is embedded in the `Chainer.py` and `Auditor.py` files to demonstrate that this new automated audit technique works as expected.

3. POSTPONED WORK

This section contains work that was omitted from phase 2 due to a lack of time to complete them. However, we include them for completeness since they are essential for future deployments of ABLS.

3.1 Client Authentication

There are two types of clients that can interact with an ABLS instance at runtime: the clients generating log data and the clients requesting log data for the purpose of auditing. To support log message confidentiality, clients generating log data will be authenticated using two-way SSL authentication, as shown in Figure 3. This will promote a bidirectional measure of trust between the client and server while at the same time encrypting all sensitive data as it is sent between the two endpoints. Also, without an appropriate certificate authority, all of the certificates will be self-signed for verification purposes. Actual deployments will require these certificates to be signed by a legitimate CA.
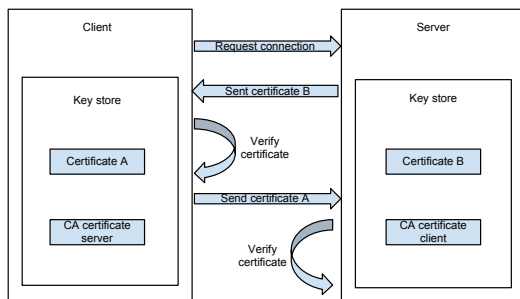

Figure 3: A pictorial representation of two-way SSL authentication.

4. ATTACK STRATEGY

My attack strategy for the next phase of the project is two-fold. First, I will use a variety of applicable security-assessment tools to automate the detection and exploitation of common vulnerabilities in web applications. Specifically, I will use IBM Appscan [1], Nessus [2], and metasploit [3] to dynamically test the behavior of the application once it is deployed. AppScan will serve to automatically scan the application to identify known vulnerabilities like SQL injection, cross-site scripting, and cross-site request forgeries. Nessus will be used analyze the network configuration of the

application deployment environment. Finally, metasploit, an open-source alternative to Nessus and Appscan, will be used in the event that neither of these two enterprise products provides sufficient results. Altogether, these tests will be dynamic in nature and rely on the application being deployed at a publicly accessible server. The goal of these tests is to determine the presence of vulnerabilities like SQL-injection, cross-site scripting, and cross-site request forgery in the web application and how they can be used to exploit the backend database.

In addition to testing the behavior of the application using these tools, I will also conduct a source-code review to determine how the database is being used. In particular, I will examine the implementation of authentication code (i.e. password hashing and persistence techniques), user input handling on both the client- and server-sides of the environment, and search for hard-coded strings embedded in the source code that reveal sensitive information.

Using the results from these tests, I will then focus intimately on the database using the guidance set forth by Ben Natan [4]. In this part of the "attack" phase I make the assumption that I will be given access to the server on which the application is deployed, so as to avoid the need to "hack" in to conduct a further analysis. Depending on the type of DBMS deployed in the application, I will follow the database hardening checklists provided by Ben Natan to verify that the databases environments are properly configured.

To finish my analysis I will examine the system configuration and source code with regard to the following elements.

1. SQL table creation, role definitions, and trigger code.

2. DB password storage locations and management techniques.

3. Audit support and configuration.

This attack strategy will cover network, application, and database vulnerabilities from a configuration and implementation perspective, and should lead to the production of a solid vulnerability assessment report.

5. REFERENCES

[1] IBM software - IBM security AppScan family - united states. http://www-01.ibm.com/software/awdtools/appscan/. Accessed: 2013-01-23.

[2] Nessus vulnerability scanner. http://www.tenable.com/products/nessus. Accessed: 2013-01-23.

[3] Penetration testing software | metasploit. http://www.metasploit.com/. Accessed: 2013-01-23.

[4] R. Ben-Natan. *Implementing database security and auditing: a guide for DBAs, information security administrators and auditors*. Digital Press, 2005.

[5] J. Bethencourt, A. Sahai, and B. Waters. Ciphertext-policy attribute-based encryption. In *Proceedings of the 2007 IEEE Symposium on Security and Privacy*, SP '07, pages 321–334, Washington, DC, USA, 2007. IEEE Computer Society.

[6] R. A. Popa, C. M. S. Redfield, N. Zeldovich, and H. Balakrishnan. Cryptdb: processing queries on an encrypted database. *Commun. ACM*, 55(9):103–111, Sept. 2012.

[7] R. Sandhu, D. Ferraiolo, and R. Kuhn. The nist model for role-based access control: towards a unified standard. In *Symposium on Access Control Models and Technologies: Proceedings of the fifth ACM workshop on Role-based access control*, volume 26, pages 47–63, 2000.